Nithesh Kumar Chavan

Charlotte, NC | +1 (336) 405-0789 | chavan3888@gmail.com | LinkedIn

SUMMARY

Software Developer with over 4 years of experience delivering high-throughput microservices, real-time Kafka pipelines, and AI-integrated systems in compliance-driven environments. Specialized in bridging scalable Java backends with modern RAG (Retrieval-Augmented Generation) architectures to turn complex datasets into actionable insights. Expert in automating enterprise CI/CD workflows using GitHub Actions and Kubernetes to accelerate release velocity. Proven track record in optimizing API response times by 35% and streamlining healthcare and fintech data retrieval using Python-based AI agents and interactive Angular frontends.

TECHNICAL SKILLS

Languages: Java, Python, TypeScript, JavaScript, SQL, PL/SQL, HTML5, CSS3
AI & Frameworks: RAG, LangChain, Vector DBs (Pinecone, Milvus), Spring Boot, FastAPI, Angular 14+, Spring Security
Cloud & DevOps: AWS (EC2, S3, EKS), GCP, Docker, Kubernetes, GitHub Actions, Terraform, Jenkins, Helm
Data & Messaging: Apache Kafka, PostgreSQL, Redis, MongoDB, Spring Data JPA, Oracle, RabbitMQ
Observability: Datadog (including AI/LLM monitoring), ELK Stack, Prometheus, Grafana, Splunk, OpenTelemetry
Tools: Maven, Gradle, Git, JIRA, Confluence, Postman, Proxyman, Xcode, SonarQube, Swagger/OpenAPI

PROFESSIONAL EXPERIENCE

Software Developer | Cotiviti — Mar 2025 – Present

- Spearheaded the development of a RAG-based analytics engine using Python and LangChain to automate the retrieval of healthcare policy data; implemented semantic search over 10,000 compliance documents which reduced manual audit research time by 60%.
- Architected a centralized monitoring dashboard in Angular 17 to visualize AI model performance; utilized RxJS and WebSockets to stream real-time metrics on token usage, inference latency, and model drift directly to the executive team.
- Engineered a robust CI/CD pipeline using GitHub Actions to automate the containerization of Java and Python services; orchestrated seamless deployments to AWS EKS clusters which eliminated manual intervention and cut release lead times by 40%.
- Refactored legacy Spring Boot microservices to move compute-intensive healthcare logic to asynchronous background workers; utilized Redis for distributed caching to achieve a sustained 35% improvement in API response times.
- Optimized Google Cloud SQL performance by identifying and refactoring expensive joins in production schemas; leveraged AI-driven performance insights to reduce complex monthly report generation time from hours to minutes.
- Integrated Datadog LLM Monitoring to track the health of production AI agents, setting up automated alerts for hallucination thresholds and ensuring high reliability for external-facing chatbot features.

Software Developer | LendingTree — Feb 2024 – Mar 2025

- Architected and deployed a resilient Java/Kafka streaming service to handle a burst load of 5,000+ financial records per 30 seconds, ensuring zero data loss and high availability across distributed broker clusters during peak mortgage application cycles.
- Designed and built modular, accessible UI components in Angular for a centralized loan tracking portal; implemented reusable form validation and reactive state management to improve the user onboarding conversion rate by 15%.
- Migrated manual infrastructure provisioning to Terraform (Infrastructure as Code), enabling the team to spin up identical staging and production environments on AWS in minutes and reducing environment-specific deployment errors by 40%.
- Secured sensitive financial APIs using Spring Security and OAuth2/JWT protocols; successfully eliminated common vulnerabilities like XSS and CSRF while passing strict quarterly security audits for client data protection.
- Implemented intelligent data validation rules using Spring AOP to flag anomalous financial records before ingestion; enhanced the overall data quality for downstream AI/ML analytics and predictive case management pipelines.
- Optimized production observability by configuring ELK stack logging and Grafana alerts for real-time system monitoring; reduced the Mean Time to Recovery (MTTR) for critical production incidents by 20% through automated error tracking.

EDUCATION

M.S., Information Technology and Management — Aug 2022 – Dec 2023
University of North Carolina at Greensboro | GPA: 3.94 — Greensboro, NC

PREVIOUS EXPERIENCE

Programmer Analyst | University of North Carolina at Greensboro — Aug 2022 – Dec 2023
Java Developer | Vadini Infocenter — Jun 2021 – Jul 2022