

Nithesh Kumar Chavan

chavan3888@gmail.com | +1 (336) 405-0789 | LinkedIn

SUMMARY

Software Engineer with 4+ years of experience designing and delivering scalable enterprise applications, distributed microservices, and high-throughput real-time data pipelines using Java, Spring Boot, Apache Kafka, and AWS/GCP. Proven ability to architect cloud-native systems from the ground up, handling everything from database optimization and API design to CI/CD automation, Kubernetes orchestration, and enterprise-grade security with OAuth2 and SAML. Experienced in fast-paced Agile environments delivering production-grade software that meets strict SLA and compliance requirements. Comfortable and experienced in AI-integrated environments, leveraging intelligent observability and building systems that support ML-powered workflows.

TECHNICAL SKILLS

Languages Java, J2EE, Python, TypeScript, JavaScript, SQL, PL/SQL, HTML5, CSS3, Node.js

Frameworks & Libraries Spring Boot, Spring MVC, Spring Data JPA, Spring Security, Spring Cloud Config, Hibernate, Angular, LangChain, FastAPI, RESTful & SOAP APIs

Messaging & Streaming Apache Kafka, RabbitMQ, Event-Driven Systems, Microservices Architecture

Cloud & DevOps AWS (EKS, EC2, S3, RDS, Lambda, CloudWatch), GCP (Cloud SQL, Cloud Datastore, Cloud Trace, Memcache), Docker, Kubernetes (EKS, AKS), Helm, Terraform, Jenkins, GitHub Actions, Ansible

AI & ML RAG Pipelines, LLM Integration (OpenAI, LangChain), Vector Databases (Pinecone, pgvector), Prompt Engineering, ML Pipeline Data Engineering

Databases PostgreSQL, Oracle, MySQL, MongoDB, Cassandra, Redis, Pinecone

Security & Observability OAuth2, SAML, JWT, Fortify, Datadog (AI/LLM Monitoring), ELK Stack, Prometheus, Grafana, Splunk, Swagger/OpenAPI

Tools Maven, Gradle, Git, JIRA, Confluence, Pandas, NumPy

PROFESSIONAL EXPERIENCE

Software Developer | Cotiviti *Charlotte, NC*

Sept 2025 – Present

- Developed scalable Spring Boot microservices and REST APIs for healthcare compliance applications handling large-scale policy and claims data processing workflows.
- Built a production RAG system using Python, LangChain, and Pinecone vector database to retrieve and rank healthcare policy documents, improving search accuracy by 60% over traditional keyword matching and reducing manual policy research time significantly.
- Implemented asynchronous event-driven architecture using Kafka and RabbitMQ for real-time claims processing, eliminating bottlenecks in data ingestion pipelines and improving end-to-end throughput.
- Developed responsive Angular 17 dashboards with RxJS observables and REST integrations to monitor AI model performance metrics, audit logs, and application health across distributed microservices.
- Containerized both Java and Python applications using Docker and orchestrated deployments to AWS EKS clusters through GitHub Actions CI/CD pipelines with automated testing and security scanning.
- Optimized API response times by 35% through strategic Redis caching implementation, SQL query tuning with explain plans, and efficient Spring Data JPA repository patterns.
- Integrated OAuth2 token-based authentication and role-based access control (RBAC) across all REST endpoints to secure sensitive healthcare data and enforce principle of least privilege.
- Designed and implemented ETL workflows in Python using Pandas and SQL to validate, transform, and load claims data from multiple sources, improving data quality metrics and enabling reliable downstream analytics.
- Monitored distributed microservices architecture using Splunk, ELK Stack, and Datadog observability platforms to identify performance bottlenecks and security anomalies in production environments.

Software Engineer | Goldman Sachs

Mar 2025 – Aug 2025

- Architected and deployed microservices on GCP and AWS using Cloud Datastore, Memcache, and Task Queues, enabling scalable low-latency communication with independent service lifecycles.
- Developed secure Spring Boot and Angular applications with OAuth2/SAML authentication, eliminating XSS/CSRF vulnerabilities and meeting enterprise security compliance standards.
- Engineered rules-driven validation and approval workflow engine that leveraged LLM-assisted policy recommendation capabilities to enable dynamic policy updates without code changes—improving audit transparency and reducing manual configuration by 40%.
- Optimized database schemas, stored procedures, and SQL queries on Google Cloud SQL; integrated AI-driven query performance analysis tools to identify slow-running queries—reducing report generation time by 50%.

- Built RESTful data ingestion APIs serving as the backend for ML model training pipelines, enabling data science teams to consume cleansed, real-time financial datasets via Spring Boot microservices.
- Enhanced observability using Google Cloud Trace and Datadog with AI/LLM model monitoring to track inference latency, token usage, and model drift—resolving bottlenecks and achieving 35% faster API response times.
- Automated CI/CD pipelines using Jenkins, Helm, Ansible, and Kubernetes for containerized microservices, accelerating release velocity and ensuring deployment consistency across environments.

Software Engineer | State of Louisiana – Dept. of Children and Family Services

Feb 2024 – Mar 2025

- Architected Spring Boot-based Apache Kafka Producer for near real-time data streaming with high availability, fault tolerance, and disaster recovery across distributed broker clusters—powering downstream AI/ML analytics and predictive case management workflows.
- Tuned Kafka performance by optimizing partition distribution, replication factors, and consumer thread pools—achieving sustained throughput of 5,000+ records per 30 seconds under production load.
- Developed RESTful APIs using Spring MVC and Spring Data JPA; externalized microservice configuration via Spring Cloud Config Server, reducing environment-specific deployment errors.
- Implemented intelligent data validation rules using Spring AOP to flag anomalous records before ingestion, enhancing data quality for reporting pipelines; adopted JUnit TDD practices to maintain code reliability.
- Built Angular 7 frontend components including search, pagination, form validation, and reusable custom directives—delivering a fully responsive user interface.

Programmer Analyst | University of North Carolina at Greensboro

Aug 2022 – Dec 2023

- Conducted C programming labs for undergraduates; supported faculty in delivering IT and entrepreneurship coursework.
- Researched emerging AI, digital technology, and cloud automation trends for academic programs and conference presentations.

Java / ServiceNow Developer | Vadini Infocenter Pvt Ltd, India

Jun 2021 – Jul 2022

- Developed and customized ServiceNow modules (Incident, Problem, Change Management) using AngularJS, Bootstrap, Business Rules, and Script Includes for full-stack workflow automation.
- Integrated third-party systems via REST APIs, managing data transformation and error handling; automated cross-departmental workflows using Flow Designer and Scheduled Jobs.
- Resolved data ingestion issues by refining front-end AngularJS logic and backend Transform Map configurations; delivered documentation and training to 20+ developers and end-users.

EDUCATION

M.S., Information Technology and Management *Aug 2022 – Dec 2023*

University of North Carolina at Greensboro **GPA: 3.94**